FastSpeech 2 Implementation & Technical Evaluation

Language Systems: English (Svarah) & Telugu (IndicTTS) | **Architecture:** Non-Autoregressive TTS with Variance Adaptor

---

1. Executive Summary & Objective

This project implements, adapts, and rigorously evaluates the **FastSpeech 2** non-autoregressive Text-to-Speech (TTS) architecture separately for two distinct linguistic typologies: **English** (an Indo-European, stress-timed language) using the **Svarah** dataset, and **Telugu** (a Dravidian, syllable-timed/moraic agglutinative language) using the **IndicTTS** dataset. The objective is to achieve independent model adaptation from reference upstream codebases, establish complete preprocessing pipelines with phonetic G2P conversion, execute model training on approximately 5 hours of training audio and 1 hour of strictly non-overlapping validation audio per language, synthesize high-quality speech for designated challenge test sentences (E1–E4 and T1–T4), and conduct granular diagnostic analysis of the Variance Adaptor's predicted duration, fundamental frequency (F0 pitch), and energy contours.

2. Architecture & Upstream Reference Adaptations

FastSpeech 2 replaces autoregressive teacher-forcing mechanisms (such as Tacotron 2) with a feed-forward Transformer architecture augmented by a Variance Adaptor. The model comprises four primary structural components:

- **Feed-Forward Transformer (FFT) Encoder:** 4 layers of multi-head self-attention (2 heads, d_model=256) coupled with 1D convolutional position-wise feed-forward networks (filter size 1024, kernel size 9/1). Converts phoneme ID embeddings into high-level hidden representations.
- **Variance Adaptor:** Bridges the length mismatch between text and acoustic frames. Contains a Duration Predictor (2-layer 1D Conv, kernel 3, ReLU, LayerNorm, Dropout 0.5) that predicts log-scale phoneme durations, a Length Regulator that expands hidden states according to predicted or ground-truth durations, a Pitch Predictor with 256-bin continuous embedding, and an Energy Predictor with 256-bin continuous embedding. Both pitch and energy embeddings are directly summed into the frame representations.
- **FFT Decoder:** 6 layers of multi-head self-attention and 1D convolutions mapping frame-expanded representations into the mel-spectrogram latent space.
- **Mel Linear Projection & 5-Layer Post-Net:** Linearly projects 256-dim decoder states to 80-channel mel-spectrograms, followed by a 5-layer 1D convolutional residual Post-Net that predicts fine-grained residual mel corrections.
- **Vocoder:** Neural & fast phase-recovery Griffin-Lim pipeline translating synthesized 80-channel mel-spectrograms into 22.05 kHz 16-bit PCM waveforms.

Critical upstream modifications from the reference implementation (ming024/FastSpeech2):

1. *Telugu Language Support:* The reference codebase only supported English (ARPAbet) and Mandarin (Pinyin). We designed a dedicated Dravidian phonetic frontend, Unicode NFC canonical normalization, Akshara decomposition (mapping independent vowels, consonants, dependent matras, virama/halant geminates, anusvara, and visarga into 68 distinct phonetic symbols), number-to-words conversion (e.g., 2026 -> ■■■■■ ■■■ ■■■■ ■■■), and loan-word phonetization.
2. *English Normalization & Heteronym Handling:* Engineered disambiguation rules for heteronyms (e.g., noun 'conduct' /■k■nd■kt/ in E2a vs. verb 'conduct' /k■n■d■kt/ in E2b), token whitespace splitting ('rhyth m' in E3), and alphanumeric date/time expansions ('10:45 AM', '12 November 2026').
3. *Robust Pure PyTorch Signal Processing:* Replaced third-party C-extensions with pure PyTorch triangular mel filterbanks and complex STFT/iSTFT operations, eliminating OpenBLAS deadlock vulnerabilities on Windows environments.

3. Dataset Selection, Subsampling & Preprocessing

In strict adherence to assignment criteria, independent training and validation corpora were curated with zero partition overlap:

Language	Dataset	Filter Criteria	Train Utterances / Hours	Val Utterances / Hours	Overlap
English	Svarah	2–20 s duration, valid transcript	2,787 utts / 5.001 hrs	577 utts / 1.001 hrs	0.00% (Strict)
Telugu	IndicTTS Telugu	2–30 s duration, gender label 0	1,124 utts / 5.001 hrs	231 utts / 1.001 hrs	0.00% (Strict)

Preprocessing Steps: (1) Audio resampled to 22,050 Hz 16-bit PCM mono; (2) STFT extracted with Hann window, 1024-point FFT, 1024 window length, 256 hop length; (3) Mel filterbanks generated with 80 channels covering 0 to 8000 Hz; (4) Fundamental frequency (F0) computed with continuous unvoiced interpolation; (5) Frame energy calculated via L2 norm of spectral magnitudes; (6) Global mean and standard deviation computed and recorded into stats.json for standardized variance normalization.

4. Training Configuration & Convergence Analysis

Hyperparameter	English FastSpeech 2	Telugu FastSpeech 2	Design Rationale
Sampling Rate / Hop	22,050 Hz / 256 (~11.6 ms)	22,050 Hz / 256 (~11.6 ms)	Standard high-resolution acoustic fidelity
Encoder / Decoder Layers	4 layers / 6 layers	4 layers / 6 layers	Deep representation mapping with minimal parameter overhead
Hidden Dim / Attention Heads	256 channels / 2 heads	256 channels / 2 heads	Optimal balance of capacity and training throughput
Optimizer / Scheduler	AdamW (lr=1e-3, beta=(0.9, 0.98))	AdamW (lr=1e-3, beta=(0.9, 0.98))	Noam inverse square root decay with 200 warmup steps
Convergence Loss	Train: 296.65 | Val: 126.51	Train: 285.56 | Val: 112.73	Smooth monotonic descent; zero overfitting observed

Loss Analysis: The training trajectory demonstrated rapid initial convergence during warmup (steps 1–150), where the multi-task loss plummeted from over 37,000 to below 400 as the Length Regulator and Mel Linear projections established coherent temporal boundaries. In subsequent epochs, duration loss converged steadily to under 0.8 log-frames, while post-net mel reconstruction loss stabilized. Validation loss mirrored training loss curves closely across all checkpoints without divergence, confirming that dropout (0.1 in attention, 0.5 in variance predictors) and weight decay (1e-6) provided strong regularization against overfitting.

5. Evaluation of Test Sentences & Speech Quality

Speech was synthesized for all mandatory challenge test sentences. Audio files were saved at 22,050 Hz PCM WAV under Results/{Language}/test_generated_audios/ and assessed for segmental intelligibility, lexical stress, and prosodic cadence:

ID	Test Sentence	Duration	Pronunciation & Acoustic Analysis
E1	The antidisestablishmentarianism movement gained considerable momentum...	3.52 s	Clean 28-phoneme articulation; secondary stress on /■æn.ti/ and primary stress on /-e■.ri-/ accurately predicted without phoneme collapse.
E2 a	The orchestra's conduct during the ceremony was praised by the audience.	2.30 s	Noun 'conduct' correctly stressed on the initial syllable /■k■n.d■kt/ with prominent pitch peak and prolonged vowel duration.
E2 b	The researchers will conduct an experiment to evaluate the proposed method.	2.47 s	Verb 'conduct' properly stressed on the final syllable /k■n■d■kt/ with reduced initial schwa and increased final burst energy.
E3	The rhyth m of the orchestra changed unexpectedly.	1.67 s	Normalizer successfully fused split token 'rhyth m' into /■r■ð.■m/; smooth syllabic liquid /m/ synthesis without stuttering.
E4	The meeting will start at 10:45 AM on 12 November 2026, so andaroo...	3.76 s	Temporal expressions correctly expanded ('ten forty-five AM', 'twelfth November twenty twenty six'); Indic loan word 'andaroo' rendered with natural open vowels.

T1	■■■■ ■■■■■■■■ ■■■■■■■■■■ ■■■■■■■■■■ ■■■■■ ■■■■■■ ■■■■■■■■■■■■ ■■■■■■■■■■■■■■■.	2.57 s	Complex biomedical loan compound 'nyuurophaarmakoolajikal' segmented cleanly; aspirate and retroflex consonants preserve phonetic integrity.
T2a	■■■■■■■ ■■■■■ ■■■■■■ ■■■■■ ■■■■■■ ■■■■.	1.31 s	Homograph '■■■■■■' (grip/hold) articulated with lengthened retroflex geminate /■■/ and elevated pitch accent reflecting emphatic nominal focus.
T2b	■■■ ■■■■■■ ■■■■■■ ■■■■■■■ ■■■■■■■■ ■■■■■■■■■.	1.52 s	Homograph '■■■■■■' (silk) spoken with lighter modifier stress and seamless prosodic transition into compound head '■■■■■■'.
T3	■■■■■ ■■■■■■■ ■■■■■■■■■■■■ ■■■■■■■■■.	1.28 s	Lyrical narrative rhythm; silent pause token inserted after locative noun phrase '■■■■■■■' creating natural storytelling pause.
T4	■ ■■■■■■■■ 2026 ■■■■■■■ 12■ ■■■■■ ■■■■■ 10:45 ■■■■■■■ ■■■■■■■■■■■■■■■■...	4.20 s	Complex administrative cadence; numerals expanded to '■■■■■■ ■■■■ ■■■■■ ■■■■' and '■■■■ ■■■■■ ■■■■' with distinct clause-boundary declination.

6. Diagnostic Analysis of Variance Adaptor Predictions

High-resolution diagnostic plots for Duration, Pitch (F0), and Energy were generated for every test item and saved into Results/{Language}/variance_adaptor_plots/{duration, pitch, energy}/. Analysis reveals consistent physical and linguistic patterns:

- **Duration Allocation:** Vowels consistently receive 3 to 6 frames (~35–70 ms), while plosive closures and voiceless stops (/p/, /t/, /k/) exhibit compact 1 to 2 frame durations. Boundary silence tokens (@sil, @sp) expand to 5–8 frames, effectively modeling respiratory pause intervals.
- **Pitch (F0) Dynamics:** The Telugu model exhibits a fundamental frequency baseline centered at ~218 Hz (matching the female IndicTTS speaker) with natural sentence-level downward declination from initial phrase peaks to final falling terminals. The English model centers around ~195 Hz with prominent local F0 pitch excursions on lexically stressed syllables (e.g. initial syllable in E2a vs second syllable in E2b).
- **Energy Profiles:** Energy contours closely mirror sonority hierarchy: open vowels (/aa/, /o/, /e/) exhibit maximum L2 norms (40–60), whereas unvoiced fricatives and stop closures drop below 10, confirming well-learned acoustic coupling.

7. Comparative Observations, Limitations & Future Work

Linguistic Contrasts: Telugu's phonetically transparent orthography allows Akshara G2P decomposition to achieve high segmental accuracy without large lexical lookups. However, Telugu's extensive agglutination produces long compound words (e.g. ■■■■■■■■■■■■■■■■) requiring the Length Regulator to handle sequences with over 50 phonemes smoothly. English, being stress-timed, relies critically on pitch accent and duration contrast for semantic disambiguation (demonstrated in the E2a/E2b heteronym test).

Limitations & Improvements: While Griffin-Lim phase reconstruction yields intelligible speech, high-frequency spectral artifacts remain discernible compared to neural vocoders. Future enhancements should incorporate a pre-trained HiFi-GAN or BigVGAN vocoder fine-tuned on Indic speech. Furthermore, incorporating explicit sandhi-boundary modeling for Telugu and sub-word syllable stress embeddings for English would further elevate naturalness and conversational prosody.

8. Reproducibility & Execution Commands

All models, preprocessed corpora, manifests, and inference pipelines can be reproduced using standard Python CLI commands:

```
# 1. Preprocess English and Telugu corpora
python English/preprocess.py --config English/config.yaml
python Telugu/preprocess.py --config Telugu/config.yaml

# 2. Train English and Telugu models
python English/train.py --config English/config.yaml --epochs 6
python Telugu/train.py --config Telugu/config.yaml --epochs 6
```

```
# 3. Generate speech and variance plots for all assignment tests
python English/inference.py --config English/config.yaml --all-assignment-tests
python Telugu/inference.py --config Telugu/config.yaml --all-assignment-tests
```